

Data Card: Atbara

This Data Card outlines the dataset composition, labeling taxonomy, web ethics, bias mitigation, and core values for the **Atbara** Pan-African financial and e-commerce complaint routing model.

Category	Template Questions	Response
Input vs. Output	• What are users inputting (for example, audio, text, images)? • What are the outputs or labels your model will predict or generate?	• Inputs: Multi-channel text messages and customer complaints spanning cross-border merchant checkouts, peer-to-peer transfers, and multi-currency marketplace purchases. • Outputs: Predicted classification labels corresponding to the 10-category operational taxonomy, paired with associated complaint priority levels (High, Medium, Low) to facilitate targeted back-office routing.
Classification problem	• What is the classification task, if any? • What categories or classes are you working with, and why?	• Task: Multi-class text classification for automated complaint routing. • Categories: Categories were pre-selected as part of the competition: 1. <i>Billing</i> (Medium) 2. <i>Account_access</i> (Low) 3. <i>Fraud_unauthorized</i> (High) 4. <i>Customer_service</i> (Medium) 5. <i>General_inquiry</i> (Low) 6. <i>Delivery_shipping</i> (High) 7. <i>Product_defect</i> (Medium) 8. <i>Refund_return</i> (High) 9. <i>Warranty_repair</i> (Medium) 10. <i>Subscription_cancel</i> (Medium)

		● Reasoning: Standard English models and rigid chatbots consistently fail to handle regional payment rails (NIP, M-Pesa, +PayShap) and localized code-switching. Categorizing complaints across these target labels allows issues to be routed directly to specialized back-office teams, reducing backlogs and cart abandonment.
Labeling	● Where do your labels come from—domain experts, community members, crowd workers, AI-generated sources? ● How will you check for consistency and accuracy?	● Sources: The labels were assigned by a Large Language Model ● Consistency and Accuracy: The Labeling guide provided by the competition organizers were passed on to an LLM model (GLM-5.2) to classify the 14,285 rows of real-world complaints scraped by the team. 947 rows were marked as unclear by the LLM, while the rest fit into the 10 given label categories. The team performed a manual check of the first 200 rows to check for accuracy of the labeling.
Data creation	● How will you collect your data? ● Will it come from direct sources (like interviews or recordings), or will you synthesize data in low-resource settings (e.g., for under-resourced languages)?	● Primary Collection: 14,285 rows of real-world, anonymized customer complaints from 4 African countries; Nigeria, Ghana, Kenya, South Africa, were scraped directly from Google Play Store. These are reviews of leading e-commerce, quick-commerce, and delivery apps including Shein, Jumia, Kilimall, Takealot, Temu, Glovo, Bolt Food, and Checkers Sixty60. ● Data Synthesis: Of the 14,285

		rows, only 52 were found to be related to warranty_repairs. The team generated 214 rows of synthetic “warranty_repairs” complaints to augment, bringing the total dataset to 14,499. The model, Gemini 2.5-flash was used for the data generation.
Web collection ethics	• Do you plan to source data from the web? • If yes, which sites or platforms? • What tools will you use - AI tools, human workers, or both? • Have you addressed issues of consent, copyright, and privacy?	• Web Sourcing: Yes, publicly accessible review channels on the Google Play Store for target Pan-African applications. • Tools: Web scraping scripts paired with human-in-the-loop validation and filtering. • Ethics & Privacy: Privacy, consent, and intellectual property considerations were systematically addressed during dataset preparation. Privacy was safeguarded by systematically filtering and scrubbing all Sensitive and Personally Identifiable Information (S/PII)—including user names, account details, phone numbers, and transaction IDs—prior to training. While user consent is inherently silent in publicly scraped data, copyright compliance was maintained by restricting collection to publicly available, non-copyright-protected review text, following prior verification of platform Terms of Service.
Bias and representation	• Who might be left out of your dataset? • How will you ensure	• Excluded Demographics: Non-digital banking clients, users using offline/USSD channels

	that the dataset reflects the diversity of the community and includes underrepresented voices?	exclusively, or customers not posting public app reviews. ● Representation & Diversity: To capture regional diversity, the dataset deliberately spans 30+ major financial and e-commerce platforms across Western, Eastern, and Southern Africa. It explicitly incorporates regional code-switching and localized dialects, including Nigerian Pidgin, Swahili, and Kenyan Sheng, alongside Standard English.
Values and community alignment	● How does your dataset design reflect the values you identified in <i>01 Build Your Own Small Language Model</i>? ● How will your community benefit from this dataset, and how will you ensure accountability to them?	● Values: Emphasizes linguistic inclusion and dispute resolution for localized African trade contexts ● Community Benefit & Accountability: Accelerates dispute resolution times, protects capital for everyday consumers facing transfer drops or fraudulent debits, and prevents revenue loss for regional merchants.